

RESEARCH ARTICLE

Power to the People! Assigning Popular Music Genres Based on Listeners' Perspective Using a Genre Tree Approach

Anonymous

Abstract

Despite popularity of playlists on streaming platforms, music genres remain a primary organizational tool among listeners, industry stakeholders, and researchers. Categorical genre assignments underpin music genre recognition in MIR and recommender systems research as well as foundational research on music preferences and functions. Yet 'correct' genre labels for popular music and the structure of genre taxonomies are often contested, raising concerns about reliability, validity, and comparability of research results. Contrarily, popular music studies have long emphasized the listener-dependent nature of genre assignments – they may be considered as implicit majority voting results of a listener community. So why not make a virtue of necessity? We present a methodology to construct balanced hierarchical genre trees from folksonomy genre votes and demonstrate its feasibility for the *MusicBrainz* folksonomy and 167,762 tracks frequently listened to in Germany across the past 50 years. The resulting Community Music Genre Tree for popular music in Germany (COMGET-G) comprises 234 genres—valuable for cultural analyses, but too detailed for many applied research tasks. We address this with an algorithm that folds such trees to any desired number of categories (FOLDCAT). For tracks lacking community votes, we develop machine-learning classifiers at multiple granularity levels using openly available music features and metadata from online databases. Exemplary results indicate that our approach can assign an appropriate genre category to virtually any popular music track, representing a multiple listeners' perspective in a transparent, reproducible, and intersubjectively comprehensible way.

Keywords: music genre recognition, hierarchical taxonomy, folksonomy, network analysis, machine learning, German popular music

1. Introduction

Although playlist-driven listening has grown with streaming platforms, genre labels remain the primary way people and institutions organize popular music: e.g., listeners search and browse by genre, and playlist editors and DJs assemble genre-based playlists (Bonini and Magaouda, 2024). While some argue classificatory practices weakened since the late twentieth century (DiMaggio, 1987) and certain genre boundaries are unstable (Silver et al., 2016), empirical evidence shows listening and social signaling remain organized around genres (Airoidi et al., 2016; Lonsdale, 2021). On commercial platforms, genre assignments are embedded in metadata and pipelines—used as search filters, playlist seeds, and recommender features. Thus, platform algorithms amplify genres' role in discov-

ery and market positioning (Valverde, 2025). For researchers, genre labels serve theoretical and methodological roles: variables in preference and reception studies (anonymized authors, XXXX), inputs to recommender research (Porcaro et al., 2024), and aggregator variables for measuring diversity and mitigating hubness in high-dimensional audio and metadata spaces (Flexer and Grill, 2016; Flexer et al., 2018).

However, 'correct' genre labels for individual popular music phenomena (i.e., tracks, artists, releases, scenes, market, playlists, labels, performances etc.) are frequently contested (van Venrooij and Schmutz, 2018): the same release may be labeled "indie rock" by listeners, "alternative" by distributors, and "post-punk" by critics. Further, not only do chosen genre labels diverge, but employed taxonomies also vary strongly

17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32

in degree of differentiation. For example, the music database *Discogs* (2024) uses a compact set of 15 high-level genre tags, whereas *Spotify* exposes many thousands of genre tags (McDonald, nd). Music databases and platforms also differ in target objects (track, release, or artist) and assignment strategies (single primary genre versus multi-label tag sets).

These differences pose methodological challenges for music information retrieval (MIR) and computational musicology: inconsistent labels degrade generalizability of supervised models (Bogdanov et al., 2019), survey items are heterogeneously interpreted (Lange et al., 2025), and, on a larger scale, noncomparability between studies undermines replication and meta-analysis (Sturm and Flexer, 2023). Despite these well-known challenges, researchers often rely pragmatically on available taxonomies, without systematically engaging with the labels' actual sources or underlying theoretical assumptions (Green et al., 2024). To address this, the present paper presents a theory-based, reproducible, and transparent procedure to derive a scalable genre taxonomy for any corpus of popular music using community votes of a matching genre folksonomy. Paired with algorithmic genre recognition, this approach can (1) surface empirical regularities and disagreements in label to inform popular music studies, and (2) provide a shared, probabilistic 'currency' for cross-disciplinary work in computational musicology and MIR—so classifiers, preference studies, and sociological analyses may operate from a common, well-documented genre mapping.

1.1 Musical Genres as hierarchical discourse formation

While MIR and computational musicology often model genres as categories, popular-music studies more commonly treat them hierarchically. Describing existing taxonomies, Pachet and Cazaly (2000) identified genealogical (e.g., *rock* parent of *alternative rock*), geographical (*hip hop* supergenre of *German hip hop*), and historical-period links (*pop* supergenre of *80s pop*). Brackett (2016, p. 8) argues that hierarchy depth can extend virtually indefinitely. Hence, music enthusiasts often visualize music genres (Crauwels, 2022). However, apart from Schreiber (2015), to our knowledge, no work has yet implemented a standardized approach to develop a genre hierarchy strictly from empirical data.

The data challenge is foundational: there is little scientific consensus on which labels to select or how to relate them hierarchically. Even expert-constructed *flat* taxonomies show limited agreement (Pachet and Cazaly, 2000), and expert classifications often diverge from listener community norms (Sordo, 2008). Adopting a power-critical cultural-studies lens (du Gay et al., 2013), we read genre discourse through Foucault's (2013) notion of *discourse formation*: artists, musicologists, industry, and listeners exert unequal in-

fluence on which labels stabilize over time (Lena and Peterson, 2008). Artists often self-assign genres as strategic acts of boundary-making to gain recognition and position themselves in scenes and markets (Negus, 1999). Musicologists' assignments reflect normative and pedagogical aims (Bourdieu, 1985). Finally, industry actors' genre assignments typically reflect how the music industry strategically positions and markets artists and releases within the field which may diverge from listener perception (Valverde, 2025). We term the industry-driven labels *distribution genres*. By contrast, aggregating listener votes treats genre assignment as implicit community majority judgment, privileging the perspective of those who engage with the music over distributors and cultural authorities.

Platforms such as *RateYourMusic* (nd), *Discogs* (2024), or *MusicBrainz* (nd) provide such crowdsourced 'ground truth' on the genre of musical objects - although mostly in non-hierarchical form. This paper proposes a new methodological approach that bases genre assignments on the digital traces of the underlying democratic voting practice and derives a bottom-up *community music genre* tree hierarchy. *MusicBrainz* seems particularly suitable for this endeavour: it maintains a whitelist of >2,000 genre tags, supports tag assignment and voting at multiple levels (track/release/artist), and exposes individual votes via its open database. Additionally, we combine our popular music theory and ethical-driven approach of identifying listener perception of music genres identity and relationships with some important research-pragmatic considerations developed in the next section.

1.2 Genre taxonomy requirements for MIR and computational musicology

When empirically deriving genre taxonomies for future MIR and computational musicology research, several practical requirements apply: First, balanced genre distributions improve statistical power and machine learning performance (Japkowicz and Stephen, 2002). However, subgenres typically have fewer track assignments than their supergenres (Bogdanov et al., 2019), producing imbalance across hierarchy levels. A novel method that derives taxonomies from hierarchical ground truth should therefore optimize balance within each hierarchy level, producing tree hierarchies in which peer genres have comparable sizes regardless of depth, ideally without information loss.

Second, different research questions require different numbers of *flat genre categories*, i.e., different degrees of genre *granularity*. Granularity denotes the detail of distinction: low granularity merges many subgenres into broad classes (e.g., *rock*, *pop*, *electronic*), while high granularity preserves fine distinctions (e.g., *death metal*, *synthwave*, *UK garage*). Surveys of musical taste may need moderate to high gran-

ularity (Lange et al., 2025, e.g., 24 genres), classification models often work better with fewer, broader categories (Genussov and Cohen, 2010), and subcultural analyses can demand very fine-grained subgenre splits (Caparrini et al., 2020). Therefore, the method should extract category sets at multiple scales from a single underlying tree-hierarchy.

Third, derived category assignments at any granularity must show validity, consistency, and generalizability beyond the ground truth used to construct them (Sturm and Flexer, 2023). External validation may be realized by comparing genre assignments with those commonly used in MIR and computational musicology. Additional validity and out-of-sample generalization can be assessed by training classifiers on the new category systems at different granularities using openly available music features and metadata.

1.3 Suitable feature dimensions for community music genre recognition

Algorithmic music-genre classification is a core MIR task. Green et al. (2024) systematically review the field and critique the frequent neglect of genre theory from the social sciences and humanities when selecting features, with performance often prioritized over theory (see Merlini, 2020, for a review of genre theories). Problematically, MIR and music psychology researchers often also assume that genre can be inferred from audio features alone (Tzanetakis and Cook, 2002; Rentfrow et al., 2011), which conflicts with the sociological view we adopt: genres are “systems of orientations, expectations and conventions that circulate between industry, text and subject” (Neale, 1980, p. 19), which emphasizes that extra-musical context alongside sound matters for genre assignments. Fabbri’s synthesis (1982) combines musical style (instrumentation, rhythm, harmony, timbre) and release milieu (subcultural context, artist positioning, market placement) as determinants of genre for listeners. While MIR toolboxes extract style features from audio signals (Bogdanov et al., 2013), open music databases (e.g., *MusicBrainz* (nd), *Genius* (2009)) provide distributional metadata (label, chart performance, platform placement, lyrics).

Lyrics occupy an intermediate position as feature dimension for genre recognition, conveying both, emotional expression and stylistic content (Brand et al., 2019), and signaling release milieu via language codes and topics that signify artists’ subcultural affiliation (Carbone et al., 2024). Attention is also needed for the *mainstream pop* metagenre Fabbri describes—eclectic, commercially optimized works that defy specific milieu affiliation—since artists from niche scenes may be relabeled as *pop* after commercial success (Holt, 2007, p. 20). Thus, popularity is a relevant feature dimension. Altogether, the anonymized authors (XXXX) recently proposed four ideal-typical fea-

ture dimensions—musical style, release context, lyrical content, and popularity—to operationalize genre following Fabbri (1982); we use that feature set in the present analysis.

1.4 Research Questions

Drawing on the above theoretical and methodological foundations, this work addressed three research questions:

RQ1 May we construct an empirically grounded, theoretically valid and sufficiently balanced hierarchical genre tree from multi-tag folksonomy votes for a given popular music corpus?

RQ2 Can we aggregate the inferred tree structure to a scalable number of distinct genre categories suitable for MIR, music recommendation and statistical applications in computational musicology?

RQ3 How valid and generalizable are such community-derived hierarchical genre category assignments for empirical music research?

a) To what extent do the community-derived categories agree with existing taxonomies of comparable granularity (e.g., *Rosamerica*, *Discogs*, *Deezer*)?

b) Can algorithmic ML-classifiers based on our approach reliably predict community-derived genre categories at multiple granularities employing only openly available and theory-derived music style, release-milieu, lyrical, and popularity features?

2. RQ1: Creating a Balanced Hierarchical Genre Tree

This section describes the method and exemplary results of the novel COMGET procedure, which creates a balanced hierarchical genre tree for popular music datasets using genre co-occurrences from folksonomy tags and optional vote counts.

2.1 Exemplary Data: Popular Tracks in Germany dataset (POPTRAG) and MusicBrainz Community genre tags

For demonstration we analyzed a corpus of 254,334 unique popular tracks released and frequently listened to in Germany over the past 50 years (POPTRAG). The corpus combines four sources: 215,702 tracks from weekly *MediaControl* German single and album charts (18 Nov 1977–21 Nov 2024; reduced to 75% where linkage to *Spotify* identifiers was successful); 26,949 *Spotify* track recommendations from 126 German-market genre seeds (retrieved 29 Nov 2024; median 325 per seed, IQR 25, range 49–343); 8,801 tracks from 100 top-featured *Spotify* playlists for Germany (retrieved on 29 Nov 2024, median 90 tracks per playlist, IQR 42, range 28–750); and, finally, 11,994 unique track entries from *Germany’s Daily Top 200 Spotify* charts (01 Jan 2017–30 Nov 2024). We excluded tracks shorter than 30 seconds (853 tracks) and kept

only unique title-artist combinations, preferring earlier release dates for duplicates.

We then retrieved *MusicBrainz* genre tags for all POPTRAG tracks, using only labels from the the *MusicBrainz* whitelist (snapshot 16 Jun 2025) and excluding non-popular-music tags (*non-music*, *interview*, *comedy*, *nature sounds*, *children's music*, *spoken word*, *standup comedy*). If track-level tags were missing we used release tags, then artist tags. For denoising, we only kept genre tags occurring for at least 50 distinct artists in the corpus to ensure sufficient discursive support for any genre. Finally, we only retained tracks that received at least two community votes on tags fulfilling all of these criteria. This removed 86,613 tracks (34%), yielding a *MusicBrainz*-tagged subset of $N = 167,721$ tracks annotated with 234 distinct music genre tags spanning 1924–2024 (median 2012, IQR 16 years). Community genre assignments ranged 1–22 tags per track (median 4, IQR 3) and 2–135 votes per track (median 6, IQR 7).

2.2 Methods: Hierarchical Genre Tree Construction (COMGET)

Using the genre-tagged and denoised POPTRAG corpus, we algorithmically constructed a hierarchical genre tree from *MusicBrainz* tag co-occurrences, extending the method of Schreiber (2015). First, we built a weighted co-occurrence network where nodes are *genre* labels and directed edges $e_{A \rightarrow B}$ quantify how strongly *genre A* predicts *genre B* on the same track. For each ordered pair (A, B) , we computed a vote-weighted conditional frequency:

$$e_{A \rightarrow B} = P(B | A) \times \frac{1}{|T_{A,B}|} \sum_{t \in T_{A,B}} \frac{v_{t,B}}{v_{t,\text{total}}} \quad (1)$$

with $T_{A,B}$ the tracks tagged with both A and B , $v_{t,B}$ the votes for B on track t , and $v_{t,\text{total}}$ the total votes on t . This yields asymmetric edges ($e_{A \rightarrow B} \neq e_{B \rightarrow A}$) combining directional co-occurrence and community vote strength (e.g., if 1,000 tracks have *metal* and 600 also have *rock*, $P(\text{rock}|\text{metal}) = 0.6$; with average *rock* vote share 0.5, $e_{\text{metal} \rightarrow \text{rock}} = 0.3$, while $e_{\text{rock} \rightarrow \text{metal}}$ can differ, e.g., 0.05).

Second, to resolve reciprocal links we kept only the positive residual of the stronger direction:

$$e'_{A \rightarrow B} = (e_{A \rightarrow B} - e_{B \rightarrow A})_+ \quad (2)$$

$$e'_{B \rightarrow A} = (e_{B \rightarrow A} - e_{A \rightarrow B})_+ \quad (3)$$

(e.g., $e'_{\text{metal} \rightarrow \text{rock}} = 0.6 - 0.05 = 0.55$, $e'_{\text{rock} \rightarrow \text{metal}} = 0$), producing a unidirectional scaffold.

Third, to enforce a strict tree we assigned each node its single strongest outgoing edge as its *supergenre*, discarding all other edges. Components lacking a *supergenre* were attached to an artificial root, *POPULAR MUSIC*, ensuring full connectivity (e.g., *metal* would at-

tach to *rock* unless a stronger edge existed, and *rock* would attach to the root if it had no *supergenre*).

Finally, we balanced the tree by optimizing leaf-node populations per hierarchical level: We estimated each *genre's* size as the sum of its assignment probabilities across tracks, computed the Gini coefficient per hierarchical level, and defined overall tree imbalance as the weighted mean tree-level Ginis. We then searched for the least-imbalanced tree by iteratively removing the weakest hierarchical link of the whole tree, reattaching the resulting dangling subtree to *POPULAR MUSIC*, and recomputing imbalance again; finally, selected the iteration with minimal overall tree imbalance.

2.3 Results: Community Genre Tree for Popular Music in Germany (COMGET-G)




Figure 1: Weighted Gini coefficient (0–1) across iterations of cutting weak genre links and reconnecting dangling subtrees to *POPULAR MUSIC*.

Level	# nodes	# tracks	Gini
1	3	total: 47,846 (29%) 12,472; 16,401; and 18,973	0.136
2	119	total: 90,381 (54%) median: 263 IQR: 789 range: 10.4 to 5832	0.643
3	110	total: 29,343 (17%) median: 156 IQR: 259 range: 14.7 to 1496	0.531
4	2	total: 192 (0.1%) 55.6 and 136	0.420

Table 1: Structure of COMGET-G at different hierarchy levels; track counts are sums of genre tag assignment probabilities in POPTRAG.

The initial POPTRAG *genre* co-occurrence network comprised 234 *genre* nodes forming a single connected component with 45.14% network density, indicating high inter-tag connectivity. Median co-occurrence per

genre was 97 (IQR 63.5; range 3 for *waltz* to 229 for *pop*), showing large variation between niche and ubiquitous genre tags. After resolving reciprocal directions, each *genre* had 29.5 incoming subgenre links on average (IQR 32; min 0; max 228 for *pop rock*) and 55 outgoing supergenre links (IQR 32; min 0 for *rock*; max 105 for *dub*).

The balance optimization (Figure 1) reattached *pop* and *hip hop* directly to *POPULAR MUSIC*—*pop* had formerly been subordinate to *rock*, and *hip hop* to *pop*. However, cutting these weakest links improved overall balance and aligns with common popular music taxonomies where they function as supergenres. The three weakest remaining links are *electronic* to *pop* (0.22), *classical* to *pop* (0.27), and *jazz* to *pop* (0.29). They lacked sufficient corpus tracks to form independent supergenres at the root. The resulting *Community Music Genre Tree for popular music in Germany* (COMGET-G, interactive version in the linked *GitHub* repository) is leaf-heavy: 89.36% of genre-nodes have no subgenres. *Rock* is the largest supergenre with 63 direct subgenres. Only *rock*, *pop*, and *hip hop* remain connected directly to *POPULAR MUSIC*, forming three broad genre clusters. The tree exhibits four hierarchy levels (excluding the artificial root); see Table 1 for structural details.

3. RQ2: Creating Genre Category Systems with Scalable Granularity

This section describes procedure and exemplary results of the novel FOLDCAT tree folding algorithm that flexibly turns COMGET-type hierarchical genre trees into genre category systems with any desired level of granularity.

3.1 Methods: Folding Algorithm (FOLDCAT)

To turn any hierarchical genre tree into a scalable number of broader *genre categories* suitable for different research problems, we developed a greedy folding algorithm (FOLDCAT) with three goals: optimal *genre category* balance, a user-selectable number of *genre categories*, and a minimum *genre size* per category for robust machine-learning and statistical analysis. *Genre size* is the expected track count per *category* (useful for single-label tagging).

FOLDCAT iteratively processes leaf *genres* beginning with those deepest in the hierarchy and with largest *size*, continually moving upwards to the root. For each *genre*, *size* is computed as the sum of vote-weighted assignment probabilities across all tracks corresponding to the *genre*. If *size* meets a pragmatically-defined threshold (e.g., $t=100$ for having enough exemplars in later ML classification), the leaf becomes a *genre category*. Otherwise the leaf, its siblings, and any already-processed descendants are absorbed into its *supergenre* and their track assignments remapped upward. This *folding procedure* continues until all re-

maining *genres* meet the minimum threshold. To obtain a desired number of *genre categories*, users pick the folding solution at a local Gini minimum using a hierarchy-weighted Gini coefficient (section 2.2) from solutions for different *size* thresholds, ensuring balanced categories across levels. A detailed flowchart of FOLDCAT is provided in the Appendix. For the exemplar COMGET-G derived from POPTRAG, we grid-searched minimum *genre sizes* from 1,000 to 15,000 (step 5) and selected four local Gini minima for comparison.

Finally, *Single-label genre categories* can be assigned to tracks by majority voting: map original tags to folded *category assignments*, sum votes per category for each track, and assign the highest-vote category. Ties are broken by choosing the category with lower total assignment probability in the corpus, favoring more specific categories.

3.2 Results: Germany's Popular Music Genre Categories




Figure 2: Weighted Gini coefficient (0–1) across COMGET-G folding solutions for minimum genre sizes s from 1,000 to 15,000 in increments of 5, using the POPTRAG corpus.

For the POPTRAG corpus, the folding algorithm found multiple locally-optimal categorical solutions for COMGET-G (Figure 2): a low-granularity solution with 5 *genre categories* (min. 12,930 tracks, Gini 0.307); medium-granularity solutions with 10 and 12 *genre categories* (min. 4,100 and 3,520 tracks; Gini 0.343 and 0.342); a high-granularity solution with 25 *genre categories* (min. 1,485 tracks, Gini 0.367); and finer solutions with 32, 34, and 38 *genre categories* (min. 1,195, 1,130, 1,020 tracks; Gini 0.376, 0.369, 0.375). We focussed subsequent analyses on solutions with 5, 12, 25, and 32 *genre categories*, representing different granularities suitable for different types of research problems (Figure 3).

COMGET-G5 contains *rock*, *pop*, *hip hop*, *electronic*, and *metal* (descending by track count). In COMGET-G12, *jazz* and *schlager* appear as distinct *pop subgenre categories*, while *rock* splits into *heavy metal*, *alternative*




Figure 3: Genre category systems for POPTRAG corpus at four different granularities (5, 12, 25, 32 *genre categories*), projected onto the original COMGET-G tree hierarchy. Colors indicate additional categories introduced at each finer levels of granularity (blue = 5; red = 12; green = 25; grey = 32). Node size represents genre size in POPTRAG (sum of vote-weighted assignment probabilities).

rock, pop rock, hard rock, and indie rock. Notably, the majority of *MusicBrainz* contributors classified heavy metal as a direct subcategory of rock rather than of metal.

COMGET-G25 adds further rock subgenre categories (punk, progressive rock, folk, folk rock, classic rock, blues rock, singer-songwriter) and pop subgenre categories (soul, classical, synth-pop, r&b). A third hierarchy level emerges: death metal under metal and power metal under heavy metal. COMGET-G32 introduces house under electronic and additional rock and pop subgenre categories (blues, country, new wave, soft rock, dance-pop, reggae). Hip hop subcategories appear only at higher granularities not shown here (pop rap at 33 categories; trap at 38).

4. RQ3: Demonstrating Validity, Reliability and Generalizability of Results

Drawing on the POPTRAG corpus of German popular music and *MusicBrainz* community genre tags, this section compares track genre assignments from our novel community categorization pipeline (COMGET + FOLDCAT) with genre category systems commonly used in MIR and computational-musicology, and reports genre-classifier training results using the same community-derived category sets. For both analyses we assigned single-label categorical genres to each track from the denoised *MusicBrainz*-tagged POPTRAG track at four granularities (COMGET-G5, COMGET-G12, COMGET-G25, COMGET-G32) using COMGET and FOLDCAT; all analyses were performed at multiple granularities.

4.1 Methods a: External Validation Strategy

For a discussion of content validity of our community genre taxonomy solutions, we employed contingency tables to compare our results with genre classifications of the same tracks in existing and frequently used taxonomies with similar granularity. For this, we retrieved *Rosamerica* genre labels from *AcousticBrainz* (Alonso-Jiménez et al., 2020), community genre assignments from the *Discogs* (2024) folksonomy, and genres assignments of the commercial music streaming platform *Deezer* (nd) for each track using our own anonymized music-data linkage tool (XXXX). We also computed category prevalences and overall Gini coefficients per taxonomy to compare genre balance. For clarity in figures, categories comprising < 2.5% of corpus tracks were excluded from visualization.

4.2 Methods b: Genre Classification of Unseen Tracks

To evaluate how well an algorithmic classifier predict the COMGET-G genre categories from FOLDCAT for unseen tracks, we trained independent classifiers for each of the four solutions. To that end, we augmented POPTRAG with 192 descriptors of musical style, release context, lyrical content, and popularity using XXXX and matched records via ISRC or fuzzy string matching

against *Spotify*, *MusicBrainz*, *AcousticBrainz*, *Genius*, and *Deezer* (API or scraping). Specifically,

retrieved features include *AcousticBrainz* high-level audio features; artist demographics (*persona*, *origin*, *birth year*, *gender*) from *MusicBrainz*; popularity metrics from *Spotify* and *Deezer*; and *Echo-Nest* high-level audio features from *Spotify*. For lyrics, we determined the language and textual statistics (*distinct words ratio* and *repeated lines ratio*) and calculated lexicon-based sentiment and emotion scores using a bag-of-words approach (Rinker, 2021; Jockers, 2015; Mohammad and Turney, 2013). The full feature list appears in Table C.1 (Appendix). To include the theorized influence of industry genre marketing on listener categorization, we derived *distribution genres* by applying COMGET + FOLDCAT to artist-level *Spotify* tags and included 27 resulting *industry genre categories* as additional predictors.

4.2.1 Learner Choice, Hyperparameter Tuning and Performance Evaluation

For a detailed ML workflow see Appendix B. We initially compared four different supervised learners using a coarse hyperparameter search with 5-fold cross-validation on a 20% stratified subsample of the training data at low (5) and medium (12) granularity to identify the most suitable model from: elastic net (GLMNET; Tay et al., 2023), regularized discriminant analysis (RDA; Weihs et al., 2005), random forest (RF; Wright and Ziegler, 2017), and light gradient boosted machine (LightGBM; Shi et al., 2025). For each we tuned key hyperparameters: GLMNET — *penalty* and *mixture* in $[0, 1]$; RDA — *covariance fraction* and *identity fraction* in $[0, 1]$; RF — *feature fraction considered at each split* in $[1, \# \text{ predictors}]$ and *leaf node size* in $[1, 100]$; number of trees was set to 500 for tuning and 1000 for final models; LightGBM — *boosting rounds* in $[50, 2000]$, *learning rate* in $[0.001, 0.1]$, *leaf node size* in $[1, 100]$, *tree depth* in $[3, 15]$, and *feature fraction* in $[1, \# \text{ predictors}]$.

To address class imbalance, we used a tunable adaptive balancing that combined down- and upsampling towards a *target* that was based on class distribution. For all learners, a *target ratio* was tuned in $[1, \max(\text{observations per class})/\min(\text{observations per class})]$. Higher granularities (25, 32) were excluded from initial learner selection due to sparse outcome levels in the 20% subsample. After choosing the best learner, we ran a more extensive hyperparameter search with 5-fold crossvalidation on the full training set for all four levels of genre granularity (5, 12, 25, 32).

Finally, tuned classifiers were trained on the full training set and evaluated on the holdout set; we also calculated model-based feature importances and confusion matrices for interpreting classification outcomes. We compared classifier performance to *Com-*

munity certainty (overall and per category), since models logically cannot exceed the consistency of human assignments (Sturm and Flexer, 2023). *Certainty* was computed as the mean assignment probability across all *MusicBrainz* votes for each track mapped to the *single-label genre* during the FOLDCAT procedure. All random operations used seed 42.

4.3 Results a: Comparison of COMGET-G variants with existing genre taxonomies

The *Rosamerica* genre classification system comprises eight categories, so we compare it to COMGET-G5; however, *AcousticBrainz* lacked entries for 28% of POPTRAG tracks, introducing a modest popularity bias into the analysis. The contingency table (Figure 4a) shows the strongest alignment between *Rosamerica*’s *Rock* and COMGET-G5’s *rock* (70%); *Rock* also is the only *Rosamerica* genre substantially associated with *metal*. Every *Rosamerica* genres is split across at least two COMGET-G5 genres ($\geq 10\%$): *Pop* into *rock* and *pop*; *Dance* into *electronic*, *pop*, and *rock*; *Jazz* into *pop* and *rock*; *Classical* into *rock*, *pop*, and *electronic*. *Rhythm and Blues* and *Hiphop* are even dispersed; *Hiphop* aligns only 37% with COMGET-G5’s *hip hop*. However, given COMGET-G5’s broad categories and *Rosamerica*’s purely feature-based method (which may be theoretically flawed), these mappings are largely reasonable. COMGET-G5 is marginally more balanced on POPTRAG ($Gini_{COMGET-G5} = 0.487$ vs. $Gini_{Rosamerica} = 0.497$).

Discogs’ 15-genre taxonomy approximates COMGET-G12 in scale. About 5% of tracks with COMGET-G labels had no *Discogs* counterpart, causing minimal popularity bias. The contingency table (Figure 4b) reveals strong alignment ($> 75\%$) for *Discogs*’ *Funk / Soul*, *Latin*, *Reggae*, *Classical*, *Hip Hop*, and *Funk*, *World & Country*. *Jazz*, *Pop*, *Rock*, and *Electronic* map most strongly to similarly named COMGET-12 categories (48%–62%). *Jazz* is also frequently associated ($\geq 10\%$) with *pop* and *rock*. Besides *pop*, *Discogs*’ *Pop* also covers the *schlager* category, a genre typical for German popular music (Thiesen, 2023). *Rock* spreads ($\geq 2\%$) across *pop* and seven COMGET-12 genres that have *rock* or *metal* in their name, suggesting *Discogs* lacks the granular rock subgenre labels for adequate German popular music representation. *Electronic* shows substantial association ($\geq 10\%$) with COMGET-12’s *pop* and *rock*. Five *Discogs* categories—*Blues*, *Brass & Military*, *Children’s*, *Non-music*, and *Stage & Screen*—each account for $< 0.5\%$ of POPTRAG, offering little discriminative value. *Folk*, *World*, & *Country*, *Classical*, *Latin*, and *Reggae* are weakly populated ($\leq 1\%$ of tracks), limiting suitability for POPTRAG. COMGET-G12 has a much more balanced distribution than *Discogs* ($Gini_{COMGET-G12} = 0.542$ vs. $Gini_{Discogs} = 0.830$).

Using the *Deezer* taxonomy, 9% of POPTRAG tracks




Figure 4: Contingency tables comparing COMGET-G genre assignments folded with granularity 5 and 12 to corresponding genre classification by taxonomies commonly used in MIR and computational musicology; values show $P(\text{COMGET} \mid \text{External})$. Additionally, we calculated relative genre prevalences and Gini coefficient across genres. Categories with less than 2.5% relative frequency in the corpus were discarded (*Rosamerica*: *Speech*, *Discogs*: *Blues*, *Brass & Military*, *Children’s*, *Non-Music*, *Stage & Screen*).

had no assignment; 27 *Deezer* genres appear in the corpus. Comparing to COMGET-G25, the contingency table (Appendix Figure D.1) shows > 75% alignment for *Rap/Hip Hop*, *Country*, *Jazz*, *Reggae*, *Dance*, *Electro*, and *Schlager & Volksmusik*. *Rock* and *Heavy Metal* distribute across COMGET-G25’s *rock* and *metal* genres. *Alternative* additionally shows overlap with *pop* and *electronic*. *Pop* aligns mostly with *pop* and *rock*, but also weakly with many genres. *R&B* overlaps most with *pop*, *soul*, *r&b*, and *hip hop*. Overall, these overlaps appear theoretically reasonable. COMGET-G25 is substantially more balanced than *Deezer* for POPTRAG ($\text{Gini}_{\text{COMGET-G25}} = 0.602$ vs. $\text{Gini}_{\text{Deezer}} = 0.822$). Fifteen out of 27 *Deezer* categories each account for < 0.5% of tracks; *Country* and *Reggae* < 1%.

4.4 Results b: Classification Performance for four different COMGET granularities

Comparison for low and medium granularity categories (Figure 5) shows all four machine-learning algorithms perform similarly on the 20% subsample (mean $F1_{\text{macro}}$; $\text{CV} = 0.717\text{--}0.758$ for COMGET-G5 and $0.468\text{--}0.503$ for COMGET-G12). LightGBM is the best performer at both granularities and was selected for all subsequent training. GLMNET ranked second for COMGET-G5, but GLMNET and RDA underperformed relative to tree-based methods on COMGET-G12, suggesting models with mainly linear or low-order quadratic boundaries struggle with finer outcome complexity.

For POPTRAG with COMGET-G5, LightGBM (hy-




Figure 5: Performance comparison of different learners (GLMNET, RDA, RF, GBM) on 20% of the data set for low and medium genre granularity; plots show mean and standard error of $F1_{\text{macro}}$ across CV folds.




Figure 6: Confusion matrices for the final GBM models on the holdout set. Values show $P(\text{Predicted}|\text{Actual})$; certainty indicates the mean assignment probability for each *genre category* across community votes.

perparameters in Appendix Table D.1) achieved $F1_{\text{macro; holdout}} = 0.743$. Per-class $F1$ correlates strongly with *community assignment certainty* ($\text{Cor}_{\text{Kendall}} = 0.949$). Given average *community certainty* of 90%, classifier performance is satisfactory. Cohen's $\kappa_{\text{holdout}} = 0.674$ on the holdout set approximates established inter-rater agreement standards ($\kappa = 0.7$) in social sciences (Hildebrandt et al., 2015, p. 327). The confusion matrix (Figure 6a) shows most errors for *electronic* (56% TP), often confused with *pop* and *rock*; *metal* (61% TP) is frequently misclassified as *rock*.

For COMGET-G12 we achieved $F1_{\text{macro; holdout}} = 0.544$ ($\kappa_{\text{holdout}} = 0.551$), which parallels lower *community certainty* (67%), indicating greater classification difficulty for both, listener community and algorithmic classifier. Per-class $F1$ and certainty remain correlated (Kendall $\tau = 0.721$). Confusions concentrate among rock subgenres (*alternative rock*, *hard rock*, *indie rock*, *pop rock*). *Heavy metal* (56% TP) is often confused with *metal* (but less often vice versa) and sometimes with *hard rock*, despite being recognized as distinct from *rock*. *indie rock* (49% TP) and *jazz* (61% TP) are frequently misclassified as *electronic*; *jazz* also mixes with *pop*. *pop rock* performs worst ($F1 = 0.182$), matching its low *community certainty* (50%), often confused with high-certainty genres *rock* and *pop*. Conversely, *schlager* and *hip hop* are classified well ($F1 = 0.837$ and 0.858 , respectively) despite differing *community certainties* (67% and 100%), indicating distinct feature signatures.

For COMGET-G25, performance further drops to $F1_{\text{macro; holdout}} = 0.348$, $\kappa_{\text{holdout}} = 0.430$, alongside lower *community certainty* (56%). The certainty- $F1$ correlation weakens (Kendall $\tau = 0.589$) but remains substantial, implying additional factors degrade performance at higher granularity. The confusion matrix (Appendix D.2a) again shows rock subgenre and metal-related confusions and a secondary confusion hub around *pop* misclassifying with multiple genres. The model performs well ($F1 > 0.6$) on distinct genres with high *certainty* (*classical*, *electronic*, *hip hop*, *metal*); poor performers ($F1 < 0.2$) are fuzzy genres with low *community certainty* (33%–50%).

For COMGET-G32 (strongly decreased *community certainty* of 54%), performance is $F1_{\text{macro; holdout}} = 0.332$, $\kappa_{\text{holdout}} = 0.339$ (Appendix D.2b). The certainty- $F1$ correlation declines further (Kendall $\tau = 0.560$), indicating *community uncertainty* is not the sole cause of reduced performance; confusion results resemble COMGET-G25 (see Appendix for details).

Figure 7 shows the top 20 most important features of the developed COMGET-G5 and COMGET-G12 algorithmic classifiers (COMGET-G25/G32 in Appendix Figure D.3). Similar features appear across granularities (e.g., *acousticness (Echo-Nest)*, *artist labelled “metal”*, *artist age*), but relative importance shifts. Among top features for COMGET-G5, we find




Figure 7: Top 20 most important features for automated classification at low and medium granularity. Values show relative feature importance from the final LightGBM models.

music-style features (e.g., *acousticness*, *speechiness*, *explicitness*) and artist persona metadata (e.g., solo vs. band, artist age). For COMGET-G12, *popularity* and *release year* become more important. For COMGET-G25/32, top features are mainly artist-metadata, indicating fine-grained distinctions rely more on release-milieu signals than high-level audio features. Lyrics features contribute minimally across classifiers.

5. Discussion

We contribute a novel and transparent procedure for constructing balanced hierarchical genre taxonomies and category assignments from listener-community genre discourse on open music folksonomies. Three key methodological contributions are: (RQ1) COMGET, an algorithm that builds a balanced hierarchical genre tree from tag co-occurrences; (RQ2) FOLDCAT, a folding algorithm that produces optimally balanced categories and final track assignments at user-specified granularities from such trees; and (RQ3) an empirical demonstration using a large German popular-music corpus (POPTRAG) and *MusicBrainz* votes showing that community-based genre assignments (a) match commonly used taxonomies (e.g., *Rosamerica*, *Discogs*, *Deezer*) while providing better category balance and coverage for our corpus, and (b) are learnable with supervised machine-learning models using publicly available, interpretable music descriptors derived from genre theory. We recommend this approach as a transparent, reproducible, and empirically grounded alternative to commercial or expert-curated genre taxonomies for future applied or basic research. The method generalizes to other music corpora and tag systems that expose tag co-occurrences (e.g., *Spotify*, *YouTube*, *SoundCloud*) and can be applied to tracks, artists, or releases.

5.1 Interpretation from the Perspective of Computational Musicology and Cultural Analysis

Our COMGET analysis of the POPTRAG dataset (RQ1) demonstrates the method’s ability to computationally unveil popular music genre discourse structures. *Rock*, *pop*, and *hip hop* form three primary clusters in POPTRAG, partially aligning with Silver et al. (2016) but differing in structure: *rock* is a distinct cluster here (vs. Silver et al.’s scattered “Niche”), *hip hop* shows little internal differentiation (as in Silver et al.), and *pop* is stylistically diverse (contrasting Silver et al.’s “Rock ’n Roll”). Direct comparison is limited by different units and contexts (*MySpace* artist social networks in 2007 vs. POPTRAG’s national listening corpus over five decades), but COMGET was designed exactly to enable future corpus- and context-sensitive comparisons of genre discourse (Pichl et al., 2017).

Crucially, hierarchical COMGET genre taxonomies from RQ1 primarily capture contemporary music listeners’ categorization of popular music rather than strict musicological genealogies (Pachet and Cazaly,

2000). COMGET-G does capture clear genealogical links (e.g., *rock* → *alternative rock*), but some placements appear counterintuitive only when viewed genealogically and invite popular music discourse interpretation rather than geneological ‘correction’: For example, *classical* and *jazz* appear as *pop* subgenres because POPTRAG contains accessible, crossover, and chart-oriented tracks from these fields and listeners often apply broad labels to popular music (Bogdanov et al., 2019). Similarly, *heavy metal* and *metal* appear as peer *rock* subgenres, reflecting shared 1970s roots, but also acknowledging later divergence (Gracyk, 2016).

As demonstrated, FOLDCAT (RQ2) provides scalable, coherent genre categories without unnecessary information loss, enabling consistent analyses at multiple granularities (e.g., coarse genre vs. fine-grained subcultural studies on POPTRAG) in a transparent way. This allows researchers to choose the best-suitable number of categories for comparable studies of playlist diversity using concentration indices (Porcario and Gómez Gutiérrez, 2019), individual listening patterns (Peterson, 1992), psychological links between personality and music choice (Anderson et al., 2021), and computational histories of genre emergence and chart dynamics (Petitbon and Hitchcock, 2022) from the perspective of listeners’ discourse.

From a philosophical point of view, both COMGET (RQ1) and FOLDCAT (RQ2) consequently implement power-critical concepts from popular music studies and cultural analysis (du Gay et al., 2013) into computational methods for making them applicable with large popular music corpora. Specifically, our approach employs radically democratic principles throughout: community votes determine co-occurrences, strongest links define hierarchies, majority voting assigns single genre labels. Resulting community genre taxonomies, genre categories, and genre category assignments reflect collective listener discourse rather than industry or musicological expert impositions, democratizing genre knowledge production for computational musicology. This should improve the ecological validity of music-preference research.

Finally, feature importances from classifiers of COMGET-G solutions at different granularities (RQ3) indicate how listener communities attribute genres: artist-level metadata dominate predictions across granularities, with distribution genres and artist demographics outweighing music style attributes—consistent with sociological emphasis on release milieu (Neale, 1980). Music style attributes help distinguish broad categories, while metadata and lyrics aid lower-level differentiation; track popularity becomes a stronger genre signal at finer resolutions (Holt, 2007). Disappointingly, bag-of-words and lexical lyric features performed weakly, suggesting LLM-based embeddings as a promising improve-

ment (Prevedello et al., 2024). Classification confusion patterns illuminate genre boundary issues: overlap among *rock* subgenres (*alternative*, *pop rock*, *indie*) and *electronic*’s permeability with *pop*, *jazz*, and *indie* reflect genuine and theoretically sound stylistic fuzziness, whereas clear separation (e.g., *schlager* vs. *hip hop*) occurs when release milieu and stylistic markers are distinct.

5.2 Benefits of the presented approach for MIR Research and Music Recommender Applications

Our method advances the state of the art in MIR as far as it extends Schreiber’s (2015) approach by several improvements (RQ1): we enforce optimally structured and fully connected genre hierarchies by removing weak hierarchical links and introducing a corpus-root container node. Our balancing mechanism finds the optimal solution without numeric thresholds, simplifying reproducibility and removing parameter tuning. The open-source implementation supports straightforward application to other corpora and tagging systems (demonstrated here using artist genre tags from *Spotify* for RQ3) and enables automatic updates of genre hierarchies within parameter-free processing pipelines.

Genre tree hierarchies in general, and specifically COMGET-derived hierarchies may be very complex, but FOLDCAT (RQ2) computationally reduces any hierarchical genre tree to manageable, well-balanced categories for downstream MIR and recommender tasks. The folding algorithm lets researchers and practitioners choose an appropriate granularity for their task and corpus size without the constraints of available external taxonomies and arbitrary subsets (Green et al., 2024). Hierarchical GINI-weighting suggests optimally class-balanced solutions across a range of choices (Japkowicz and Stephen, 2002), akin to hyperparameter tuning. Furthermore, full content consistency achieved by mere folding of underlying COMGET ground-truth enables to create comparable benchmarks at different levels of granularity for a target corpus. Underlying category assignments remain probabilistic and can be treated statistically, reflecting listener uncertainty.

COMGET+FOLDCAT generally align with common MIR taxonomies (RQ3) but uniquely capture categories salient to a given corpus, improving balance and coverage versus arbitrary external systems. For example, neither *Discogs* nor *Deezer* provided sufficient detail for rock subgenres in our German popular-music corpus and included irrelevant categories. Capturing corpus-specific, audience-oriented categories benefits recommender development, diversity comparisons across heterogeneous corpora, and tailoring high-level audio descriptors to listener perception (Flexer and Grill, 2016). Generally, the audience-aligned labels in our approach give a more realistic estimate of classifier generalization beyond benchmarks (Green et al., 2024), likely improving real-world recommen-

dation performance, and support building listener-centered recommender systems and diversity assessments (Deldjoo et al., 2024; Porcaro et al., 2021).

More broadly, our approach addresses a core methodological gap in MIR: most genre-recognition methods lack explicit theoretical grounding (Green et al., 2024). In contrast, we operationalize genre-as-discourse (Foucault, 2013; Neale, 1980), treating genre assignment as an empirical question rather than an ontological given. We also provide a theoretical justification for our choice of audio features, metadata, and lyric features (anonymous authors, XXXX). This links computational methods with popular music studies and humanities scholarship on cultural categorization (DiMaggio, 1987). Finally, using *MusicBrainz* as the basis for COMGET and FOLDCAT also enables to assign genre categories to virtually any popular music track, combining *MusicBrainz*’ huge scope with reliable classifiers applicable to untagged tracks.

5.3 Limitations

Several constraints affect interpretation and generalization. The *MusicBrainz* platform (launched 2000) does not provide individual vote timestamps for genre tags (*MusicBrainz*, nd). COMGET-G therefore aggregates roughly 25 years of listener votes across a period in which genre boundaries likely shifted, potentially misrepresenting historical perception for older tracks tagged much later (Nie, 2022). The tagging community may also differ systematically from the broader listening public; editor identities are withheld ($\approx 1.2M$ anonymous editor accounts), preventing assessment of sociodemographic or expertise biases. Timeline analysis shows that a substantial portion of votes stems from Bogdanov et al.’s (2019) efforts to add genre votings to the *Musicbrainz* database in 2021, sourcing the music platforms *Last.fm*, *beaTunes*, and *Discogs*.

Our tree-construction furthermore enforces strict supergenre-subgenre hierarchies with single-supergenre topology, which can oversimplify relationships where genres legitimately have multiple supergenres from the listeners’ perspective (Schreiber, 2016). This hierarchy also does not model peer-genre similarities (i.e., equal distance from common supergenre), potentially obscuring lateral relations that listeners perceive as equally salient as vertical hierarchical links. We accepted this constraint because alternative foldings would create logical inconsistencies that preclude a coherent flat categorization.

Converting probabilistic multi-tag assignments to single categorical labels for classification incurs information loss. External comparisons with *Rosamerica*, *Discogs*, and *Deezer* treat discrete mappings as equivalent despite underlying vote uncertainty. Our tie-breaking favors rarer categories to retain granularity, which may bias results toward niche genres and overrepresent subcultural boundaries less salient to

broader audiences.

Genre classification performance is below what is typically considered state-of-the-art in the field, but we intentionally used interpretable features and avoided deep learning to preserve straightforward interpretability. Given empirical limits set by community certainty and the dataset's broad, 50-year popular-music coverage, achieved performance seems reasonable. The observed strong correlation between community certainty and classifier performance supports grounding MIR evaluation in listener-perception limits rather than arbitrary absolute numeric targets.

Finally, while our approach has proven fruitful for German popular music, generalizability beyond remains unclear. Schreiber (2015) demonstrated that results for taggers using different languages on a commercial platform are broadly similar but exhibit systematic differences, implying cultural and linguistic effects on genre discourse. Future work should test COMGET + FOLDCAT across other cultural contexts, musical traditions, and listener communities beyond German and also Western (Huang et al., 2023) popular music.

5.4 Future Directions

We will pursue several complementary directions to strengthen and extend the presented framework. First, systematically comparing COMGET-G with alternative folksonomy hierarchies—including *MusicBrainz*'s own editorial hierarchy and Schreiber's (2016) solution would provide an additional validity check. Convening German popular-music experts to review COMGET-G and label prototypical tracks at multiple granularities would produce validated reference data for taste studies. Analyzing more comprehensive, less popularity-biased catalogues (e.g., the German National Library's register of music released in Germany) would extend applicability beyond very popular music and yield genre descriptors relevant across Germany's musical output. Replicating the pipeline on US and UK track corpora could reveal how national listening cultures reshape genre structures and boundaries.

Beyond academia, our approach may also offer practical value for music-industry stakeholders. Radio programmers, playlist editors, and music librarians could better understand how audiences organize the collections they choose their repertoire from. A&R teams might detect gaps between industry positioning and listener reception. Marketing departments could assess which subcultural boundaries listeners perceive in comparison to industry-imposed categories. Music journalists could ground critical analyses in empirical listener consensus. These applications demonstrate how listener-centered hierarchical genre taxonomies as the one developed here can bridge scholarly analysis and industry practice.

Competing Interests

The authors have no competing interests to declare.

References

- Airolidi, M., Beraldo, D., and Gandini, A. (2016). Follow the algorithm: An exploratory investigation of music on YouTube. *Poetics*, 57:1–13. <https://doi.org/10.1016/j.poetic.2016.05.001>.
- Alonso-Jiménez, P., Bogdanov, D., Pons, J., and Serra, X. (2020). Tensorflow audio models in essentia. In *2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 266–270.
- Anderson, I., Gil, S., Gibson, C., Wolf, S., Shapiro, W., Semerci, O., and Greenberg, D. M. (2021). “Just the Way You Are”: Linking Music Listening on Spotify and Personality. *Social Psychological and Personality Science*, 12(4):561–572.
- Bogdanov, D., Porter, A., Schreiber, H., Urbano, J., and Oramas, S. (2019). The Acousticbrainz Genre Dataset: Multi-Source Multi-Level, Multi-Label, and Large-Scale. In *20th International Society for Music Information Retrieval Conference*, Delft.
- Bogdanov, D., Wack, N., Gómez, E., Gulati, S., Herrera, P., Mayor, O., Roma, G., Salamon, J., Zapata, J., and Serra, X. (2013). ESSENTIA: an open-source library for sound and music analysis. In *Proceedings of the 21st ACM international conference on Multimedia*, pages 855–858, New York. Association for Computing Machinery.
- Bonini, T. and Magaouda, P. (2024). *Platformed! How Streaming, Algorithms and Artificial Intelligence are Shaping Music Cultures*. Springer Nature Switzerland. <https://10.1007/978-3-031-43965-0>.
- Bourdieu, P. (1985). The market of symbolic goods. *Poetics*, 14(1):13–44. [https://doi.org/10.1016/0304-422X\(85\)90003-8](https://doi.org/10.1016/0304-422X(85)90003-8).
- Brackett, D. (2016). *Categorizing Sound: Genre and Twentieth-Century Popular Music*. University of California Press. <https://doi.org/10.1525/california/9780520248717.001.0001>.
- Brand, C. O., Acerbi, A., and Mesoudi, A. (2019). Cultural evolution of emotional expression in 50 years of song lyrics. *Evolutionary Human Sciences*, 1:e11. <https://doi.org/10.1017/ehs.2019.11>.
- Caparrini, A., Arroyo, J., Pérez-Molina, L., and Sánchez-Hernández, J. (2020). Automatic sub-genre classification in an electronic dance music taxonomy. *Journal of New Music Research*, 49(3):269–284. <https://doi.org/10.1080/09298215.2020.1761399>.
- Carbone, L., Alvarez-Cueva, P., and Vandenbosch, L. (2024). Status Markers in Popular Music Across Six Countries: A Content Analysis of Gender, Race/Ethnicity, Genre, and Capital in Music Lyrics. *Sex Roles*, 90(7):891–909. <https://doi.org/10.1007/s11199-024-01483-0>.

- Crauwels, K. (2022). Musicmap. the genealogy and hierarchy of musical genres. <https://musicmap.info/>.
- Deezer (n.d.). Deezer music streaming platform. Retrieved 2025-12-12 from <https://developers.deezer.com/login?redirect=/api>.
- Deldjoo, Y., Schedl, M., and Knees, P. (2024). Content-driven music recommendation: Evolution, state of the art, and challenges. *Computer Science Review*, 51:100618. <https://doi.org/10.1016/j.cosrev.2024.100618>.
- DiMaggio, P. (1987). Classification in Art. *American Sociological Review*, 52(4):440. <https://doi.org/10.2307/2095290>.
- Discogs (2024). Database guidelines. 9. genres and styles. <https://support.discogs.com/hc/en-us/articles/360005055213-Database-Guidelines-9-Genres-Styles>.
- du Gay, P., Hall, S., Janes, L., Madsen, A. K., Mackay, H., and Negus, K. (2013). *Doing Cultural Studies: The Story of the Sony Walkman*. SAGE.
- Fabbri, F. (1982). A Theory of Musical Genres: Two Applications. In Horn, D. and Tagg, P., editors, *Popular Music Perspectives*, pages 52–81. IASPM.
- Flexer, A., Dörfler, M., Schlüter, J., and Grill, T. (2018). Hubness as a Case of Technical Algorithmic Bias in Music Recommendation. In *2018 IEEE International Conference on Data Mining Workshops (ICDMW)*. <https://doi.org/10.1109/ICDMW.2018.00154>.
- Flexer, A. and Grill, T. (2016). The Problem of Limited Inter-rater Agreement in Modelling Music Similarity. *Journal of New Music Research*, 45(3):239–251. <https://doi.org/10.1080/09298215.2016.1200631>.
- Foucault, M. (2013). *Archaeology of Knowledge*. Routledge, London, 2 edition. <https://doi.org/10.4324/9780203604168>.
- Genius (2009). About genius. <https://genius.com/Genius-about-genius-annotated>.
- Genussov, M. and Cohen, I. (2010). Musical genre classification of audio signals using geometric methods. In *18th European Signal Processing Conference*, pages 497–501. <https://ieeexplore.ieee.org/document/7096332>.
- Gracyk, T. (2016). Heavy metal: Genre? Style? Subculture? *Philosophy Compass*, 11(12):775–785. <https://doi.org/10.1111/phc3.12386>.
- Green, O., Sturm, B. L., Born, G., and Wald-Fuhrmann, M. (2024). A critical survey of research in music genre recognition. In *25th International Society for Music Information Retrieval Conference*, San Francisco. https://ismir2024program.ismir.net/poster_149.html.
- Hildebrandt, A., Jäckle, S., Wolf, F., and Heindl, A. (2015). Inhaltsanalyse. In *Methodologie, Methoden, Forschungsdesign*, pages 299–333. Springer Fachmedien Wiesbaden, Wiesbaden. https://doi.org/10.1007/978-3-531-18993-2_12.
- Holt, F. (2007). Genre in Popular Music. In *Genre in Popular Music*. University of Chicago Press. <https://doi.org/10.7208/9780226350400>.
- Huang, R. S., Holzapfel, A., Sturm, B. L. T., and Kaila, A.-K. (2023). Beyond diverse datasets: Responsible mir, interdisciplinarity, and the fractured worlds of music. *Transactions of the International Society for Music Information Retrieval*. <https://doi.org/10.5334/tismir.141>.
- Japkowicz, N. and Stephen, S. (2002). The class imbalance problem: A systematic study. *Intelligent Data Analysis*, 6(5):429–449. <https://doi.org/10.3233/IDA-2002-6504>.
- Jockers, M. L. (2015). Syuzhet: Extract sentiment and plot arcs from text. <https://github.com/mjockers/syuzhet>.
- Lange, E. B., Gernandt, E., and Merrill, J. (2025). Categorizing music by genres. *Annals of the New York Academy of Sciences*, page nias.15404. <https://doi.org/10.1111/nias.15404>.
- Lena, J. C. and Peterson, R. A. (2008). Classification as Culture: Types and Trajectories of Music Genres. *American Sociological Review*, 73(5):697–718. <https://doi.org/10.1177/000312240807300501>.
- Lonsdale, A. J. (2021). Musical taste, in-group favoritism, and social identity theory: Re-testing the predictions of the self-esteem hypothesis. *Psychology of Music*, 49(4):817–827. <https://doi.org/10.1177/0305735619899158>.
- McDonald, G. (n.d.). Every noise at once. Retrieved 2024-04-23 from <https://everynoise.com/>.
- Merlini, M. (2020). A critical (and interdisciplinary) survey of popular music genre theories. *Interdisciplinary Studies in Musicology*, 20:79–94. <https://www.cceol.com/search/article-detail?id=1105289>.
- Mohammad, S. M. and Turney, P. D. (2013). Crowdsourcing a word-emotion association lexicon. *Computational Intelligence*, 29(3):436–465. <https://doi.org/10.1111/j.1467-8640.2012.00460.x>.
- MusicBrainz (n.d.). Database statistics. Retrieved 2025-12-12 from <https://musicbrainz.org/statistics/>.
- Neale, S. (1980). *Genre*. British Film Institute.
- Negus, K. (1999). *Music Genres and Corporate Cultures*. Routledge, London. <https://doi.org/10.4324/9780203169469>.
- Nie, K. (2022). Inaccurate Prediction or Genre Evolution? Rethinking Genre Classification. *23th International Society for Music Information Retrieval Conference*. <https://doi.org/10.5281/zenodo.7316661>.
- Pachet, F. and Cazaly, D. (2000). A taxonomy of musical genres. In *Content-Based Multimedia Information Access - Volume 2*, RIAO '00, pages 1238–1245,

- Paris. <https://dl.acm.org/doi/10.5555/2856151.2856177>.
- Peterson, R. A. (1992). Understanding audience segmentation: From elite and mass to omnivore and univore. *Poetics*, 21(4):243–258. [https://doi.org/10.1016/0304-422X\(92\)90008-Q](https://doi.org/10.1016/0304-422X(92)90008-Q).
- Petitbon, A. M. and Hitchcock, D. B. (2022). What kind of music do you like? a statistical analysis of music genre popularity over time. *Journal of Data Science*, 20(2):168–187. 10.6339/22-JDS1040.
- Pichl, M., Zangerle, E., Specht, G., and Schedl, M. (2017). Mining Culture-Specific Music Listening Behavior from Social Media Data. In *IEEE International Symposium on Multimedia (ISM)*, pages 208–215.
- Porcaro, L., Castillo, C., and Gómez, E. (2021). Diversity by Design in Music Recommender Systems. *Transactions of the International Society for Music Information Retrieval*, 4(1). <https://doi.org/10.5334/tismir.106>.
- Porcaro, L. and Gómez Gutiérrez, E. (2019). 20 years of playlists: A statistical analysis on popularity and diversity. In *20th International Society for Music Information Retrieval Conference*, Deflt.
- Porcaro, L., Gómez, E., and Castillo, C. (2024). Assessing the Impact of Music Recommendation Diversity on Listeners: A Longitudinal Study. *ACM Trans. Recomm. Syst.*, 2(1):3:1–3:47. <https://doi.org/10.1145/3608487>.
- Prevedello, G., Blin, I., Monechi, B., and Ubaldi, E. (2024). Lyrics for Success: Embedding Features for Song Popularity Prediction. In Kruspe, A., Oramas, S., Epure, E. V., Sordo, M., Weck, B., Doh, S., Won, M., Manco, I., and Meseguer-Brocal, G., editors, *Proceedings of the 3rd Workshop on NLP for Music and Audio (NLP4MusA)*, pages 75–80, Oakland. Association for Computational Linguistics.
- RateYourMusic (n.d.). All genres. Retrieved 2025-12-12 from <https://rateyourmusic.com/genres/>.
- Rentfrow, P. J., Goldberg, L. R., and Levitin, D. J. (2011). The structure of musical preferences: A five-factor model. *Journal of Personality and Social Psychology*, 100(6):1139–1157. <https://doi.org/10.1037/a0022406>.
- Rinker, T. W. (2021). sentimentr: Calculate text polarity sentiment. <https://github.com/trinker/sentimentr>.
- Schreiber, H. (2015). Improving Genre Annotations for the Million Song Dataset. In *16th International Society for Music Information Retrieval Conference*, Malaga.
- Schreiber, H. (2016). Genre Ontology Learning: Comparing Curated with Crowd-Sourced Ontologies. In *17th International Society for Music Information Retrieval Conference*, New York.
- Shi, Y., Ke, G., Soukhavong, D., Lamb, J., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., Liu, T.-Y., Titov, N., and Cortes, D. (2025). *lightgbm: Light Gradient Boosting Machine*. R package version 4.6.0.99.
- Silver, D., Lee, M., and Childress, C. C. (2016). Genre Complexes in Popular Music. *PLOS ONE*, 11(5):e0155471. <https://doi.org/10.1371/journal.pone.0155471>.
- Sordo, M. (2008). The quest for musical genres: Do the experts and the wisdom of crowds agree? In *9th International Society for Music Information Retrieval Conference*, Philadelphia.
- Sturm, B. L. T. and Flexer, A. (2023). A Review of Validity and Its Relationship to Music Information Research. In *24th International Society for Music Information Retrieval Conference*, Milan.
- Tay, J. K., Narasimhan, B., and Hastie, T. (2023). Elastic net regularization paths for all generalized linear models. *Journal of Statistical Software*, 106(1):1–31. <https://10.18637/jss.v106.i01>.
- Thiesen, F. C. (2023). Taking stock of a discourse: An integrative content analysis of musical structures and cultural stereotypes in top-10 German Schlager songs from 2009 to 2019. *Popular Music*, 42(4):375–394. <https://doi.org/10.1017/S0261143024000102>.
- Tzanetakis, G. and Cook, P. (2002). Musical genre classification of audio signals. *IEEE Transactions on Speech and Audio Processing*, 10(5):293–302. <https://doi.org/10.1109/TSA.2002.800560>.
- Valverde, R. C. (2025). Inequity by design: Music streaming taxonomies as ruinous infrastructure. In Hesmondhalgh, D., editor, *Music Streaming around the World*, pages 191–210. University of California Press, 1 edition. <https://doi.org/10.2307/jj.33903674.16>.
- van Venrooij, A. and Schmutz, V. (2018). Categorical ambiguity in cultural fields: The effects of genre fuzziness in popular music. *Poetics*, 66:1–18. <https://doi.org/10.1016/j.poetic.2018.02.001>.
- Weihs, C., Ligges, U., Luebke, K., and Raabe, N. (2005). klar analyzing german business cycles. In Baier, D., Decker, R., and Schmidt-Thieme, L., editors, *Data Analysis and Decision Support*, pages 335–343, Berlin. Springer-Verlag.
- Wright, M. N. and Ziegler, A. (2017). ranger: A fast implementation of random forests for high dimensional data in C++ and R. *Journal of Statistical Software*, 77(1):1–17. 10.18637/jss.v077.i01.